

Rajesh M Venkatesan

Staff Data & AI Architect | +91 8553797755 • mvrajeshiyer@gmail.com • [linkedin.com/in/mvrajeshiyer](https://www.linkedin.com/in/mvrajeshiyer) • Bengaluru

SUMMARY

Staff Data & AI Architect with 14+ years designing large-scale data platforms and distributed systems. Proven track record of leading multi-cloud migrations, building high-performance data platforms, and delivering measurable business impact (50% cost reduction, 3x performance improvement, 100M+ events/day pipelines). Specialize in system architecture, platform scalability, and production-grade AI systems.

TECHNICAL SKILLS

- Data Platforms:** Snowflake, BigQuery, Databricks, Spark, Airflow, dbt, Data Modeling
- Cloud & Infra:** AWS, GCP, Azure, Kubernetes, CI/CD
- AI Systems:** RAG architectures, LLM integrations, Vector DBs (FAISS, Weaviate), AI Agents(Crew AI, Langgraph), LLM's(RL, KV Cache)
- Observability:** ClickHouse, OpenTelemetry, distributed tracing
- Languages:** Python, SQL
- Leadership:** Architecture strategy, roadmap ownership, cross-team influence, mentoring

EXPERIENCE

- Resident Solutions Architect** | [Acceldata](#) Nov 2025 – Present
- Defined reference architecture for ADOC (Cloud Observability) and ADM (Agentic Data Management) across multi-cloud environments (AWS, GCP, Azure)
 - Standardized integration patterns, reducing enterprise onboarding time by 40%
 - Led architecture discovery workshops with Fortune 500 customers to identify high-impact use cases and define scalable solutions
 - Influenced product roadmap by translating field insights into platform improvements
 - Designed scalable observability and pipeline monitoring architectures across complex data ecosystems
- Data Architect (Product)** | [VuNet Systems](#) May 2024 – Nov 2025
- Led performance architecture and benchmarking for ClickHouse, Kafka, and OpenTelemetry pipelines
 - Improved dashboard performance by 30–50% using query optimization, materialized views, and tiered storage
 - Reduced system incidents by 70% through architectural improvements and data retention strategies
 - Defined best practices for data lifecycle management, storage optimization, and observability design
 - Solved large-scale system bottlenecks across CPU, memory, and storage layers in production environments
- Engineering Manager** | [Quotient Technology Inc.](#) Jun 2022 – May 2024
- Owned architecture for 1PB data platform migration (on-prem → Snowflake/GCP)
 - Designed scalable data systems handling 2x YoY growth across analytics and ML workloads
 - Optimized Airflow pipelines processing 10TB/day across 200+ workflows
 - Built internal deployment framework reducing release cycles by 40%
 - Led team of 8 engineers while maintaining hands-on architecture ownership
- Lead / Senior / Data Engineer** | [Quotient Technology Inc.](#) Mar 2016 – Jun 2022
- Progressed from IC → Lead while owning core data platform components
 - Built and scaled pipelines supporting analytics and ML systems under rapid data growth
 - Contributed to architecture evolution of data platform during scale-up phase

EARLIER EXPERIENCE

- Software Engineer** | CenturyLink May 2014– Jan 2016
- Application Developer** | IBM Aug 2011 – May 2014

EDUCATION

- M.Tech Computational Engineering and Networking** | Amrita Vishwa Vidyapeetham 2009 – 2011

AWARDS

* Performance Award from IBM Client * Mountain Mover Award from Quotient Technology